

```
In [1]: from dsc80_utils import *
```

Lecture 5 – Exploratory Data Analysis and Data Cleaning

DSC 80, Fall 2025

Agenda

- Merging datasets.
- Dataset overview.
- Introduction to `plotly`.
- Exploratory data analysis and feature types.
- Data cleaning.
 - Data quality checks.
 - Missing values.
 - Transformations and timestamps.
 - Modifying structure.
- Investigating student-submitted questions!

Merging

Example: Name categories

The [New York Times article from Lecture 1](#) claims that certain categories of names are becoming more popular. For example:

- Forbidden names like Lucifer, Lilith, Kali, and Danger.
- Evangelical names like Amen, Savior, Canaan, and Creed.
- Mythological names.
- It also claims that baby boomer names are becoming less popular.

Let's see if we can verify these claims using data!

Loading in the data

Our first DataFrame, `baby`, is the same as we saw in Lecture 1. It has one row for every combination of `'Name'`, `'Sex'`, and `'Year'`.

```
In [2]: baby_path = Path('data') / 'baby.csv'
        baby = pd.read_csv(baby_path)
        baby
```

```
Out[2]:
```

	Name	Sex	Count	Year
0	Liam	M	20456	2022
1	Noah	M	18621	2022
2	Olivia	F	16573	2022
...
2085155	Wright	M	5	1880
2085156	York	M	5	1880
2085157	Zachariah	M	5	1880

2085158 rows × 4 columns

Our second DataFrame, `nyt`, contains the New York Times' categorization of each of several names, based on the aforementioned article.

```
In [3]: nyt_path = Path('data') / 'nyt_names.csv'
        nyt = pd.read_csv(nyt_path)
        nyt
```

```
Out[3]:
```

	nyt_name	category
0	Lucifer	forbidden
1	Lilith	forbidden
2	Danger	forbidden
...
20	Venus	celestial
21	Celestia	celestial
22	Skye	celestial

23 rows × 2 columns

Issue: To find the number of babies born with (for example) forbidden names each year, we need to combine information from both `baby` and `nyt`.

Merging

- We want to link rows from `baby` and `nyt` together whenever the names match up.

- This is a **merge** (`pandas` term), i.e. a **join** (SQL term).
- A merge is appropriate when we have two sources of information **about the same individuals** that is **linked by a common column(s)**.
- The common column(s) are called the **join key**.

Example merge

Let's demonstrate on a small subset of `baby` and `nyt`.

```
In [4]: nyt_small = nyt.iloc[[11, 12, 14]].reset_index(drop=True)

names_to_keep = ['Julius', 'Karen', 'Noah']
baby_small = (baby
              .query("Year == 2020 and Name in @names_to_keep")
              .reset_index(drop=True)
              )

dfs_side_by_side(baby_small, nyt_small)
```

	Name	Sex	Count	Year		nyt_name	category
0	Noah	M	18407	2020	0	Karen	boomer
1	Julius	M	966	2020	1	Julius	mythology
2	Karen	F	330	2020	2	Freya	mythology
3	Noah	F	306	2020			
4	Karen	M	6	2020			

```
In [5]: baby_small.merge(nyt_small, left_on='Name', right_on='nyt_name')
```

```
Out[5]:
```

	Name	Sex	Count	Year	nyt_name	category
0	Julius	M	966	2020	Julius	mythology
1	Karen	F	330	2020	Karen	boomer
2	Karen	M	6	2020	Karen	boomer

The `merge` method

- The `merge` DataFrame method joins two DataFrames by columns or indexes.
 - As mentioned before, "merge" is just the `pandas` word for "join."
- When using the `merge` method, the DataFrame before `merge` is the "left" DataFrame, and the DataFrame passed into `merge` is the "right" DataFrame.

- In `baby_small.merge(nyt_small)`, `baby_small` is considered the "left" DataFrame and `nyt_small` is the "right" DataFrame; the columns from the left DataFrame appear to the left of the columns from right DataFrame.
- By default:
 - If join keys are not specified, all shared columns between the two DataFrames are used.
 - The "type" of join performed is an inner join. **This is the only type of join you saw in DSC 10, but there are more, as we'll now see!**


Join types: inner joins

```
In [6]: baby_small.merge(nyt_small, left_on='Name', right_on='nyt_name')
```

```
Out[6]:
```

	Name	Sex	Count	Year	nyt_name	category
0	Julius	M	966	2020	Julius	mythology
1	Karen	F	330	2020	Karen	boomer
2	Karen	M	6	2020	Karen	boomer

- Note that `'Noah'` and `'Freya'` do not appear in the merged DataFrame.
- This is because there is:
 - no `'Noah'` in the right DataFrame (`nyt_small`), and
 - no `'Freya'` in the left DataFrame (`baby_small`).
- The default type of join that `merge` performs is an **inner join**, which keeps the **intersection** of the join keys.



Different join types

We can change the type of join performed by changing the `how` argument in `merge`. Let's experiment!

```
In [7]: # Note the NaNs!
baby_small.merge(nyt_small, left_on='Name', right_on='nyt_name', how='left')
```

Out [7]:

	Name	Sex	Count	Year	nyt_name	category
0	Noah	M	18407	2020	NaN	NaN
1	Julius	M	966	2020	Julius	mythology
2	Karen	F	330	2020	Karen	boomer
3	Noah	F	306	2020	NaN	NaN
4	Karen	M	6	2020	Karen	boomer

In [8]: `baby_small.merge(nyt_small, left_on='Name', right_on='nyt_name', how='right')`

Out [8]:

	Name	Sex	Count	Year	nyt_name	category
0	Karen	F	330.0	2020.0	Karen	boomer
1	Karen	M	6.0	2020.0	Karen	boomer
2	Julius	M	966.0	2020.0	Julius	mythology
3	NaN	NaN	NaN	NaN	Freya	mythology

In [9]: `baby_small.merge(nyt_small, left_on='Name', right_on='nyt_name', how='outer')`


Out [9]:

	Name	Sex	Count	Year	nyt_name	category
0	NaN	NaN	NaN	NaN	Freya	mythology
1	Julius	M	966.0	2020.0	Julius	mythology
2	Karen	F	330.0	2020.0	Karen	boomer
3	Karen	M	6.0	2020.0	Karen	boomer
4	Noah	M	18407.0	2020.0	NaN	NaN
5	Noah	F	306.0	2020.0	NaN	NaN

Different join types handle mismatches differently

There are four types of joins.

- **Inner:** keep **only** matching keys (intersection).
- **Outer:** keep **all** keys in both DataFrames (union).
- **Left:** keep all keys in the left DataFrame, whether or not they are in the right DataFrame.
- **Right:** keep all keys in the right DataFrame, whether or not they are in the left DataFrame.
 - Note that `a.merge(b, how='left')` contains the same information as `b.merge(a, how='right')`, just in a different order.



Notes on the `merge` method

- `merge` is flexible – you can merge using a combination of columns, or the index of the DataFrame.
- If the two DataFrames have the same column names, `pandas` will add `_x` and `_y` to the duplicated column names to avoid having columns with the same name (change these the `suffixes` argument).
- There is, in fact, a `join` method, but it's actually a wrapper around `merge` with fewer options.
- **As always, the [documentation](#) is your friend!**

Lots of `pandas` operations do an implicit outer join!

- `pandas` will almost always try to match up index values using an outer join.
- It won't tell you that it's doing an outer join, it'll just throw `NaN` s in your result!

```
In [10]: df1 = pd.DataFrame({'a': [1, 2, 3]}, index=['hello', 'dsc80', 'students'])
df2 = pd.DataFrame({'b': [10, 20, 30]}, index=['dsc80', 'is', 'awesome'])
dfs_side_by_side(df1, df2)
```

	a		b
hello	1	dsc80	10
dsc80	2	is	20
students	3	awesome	30

```
In [11]: df1['a'] + df2['b']
```

```
Out[11]: awesome      NaN
dsc80      12.0
hello      NaN
is         NaN
students   NaN
dtype: float64
```

```
In [12]: pd.concat([df1, df2])
```

Out [12]:

	a	b
hello	1.0	NaN
dsc80	2.0	NaN
students	3.0	NaN
dsc80	NaN	10.0
is	NaN	20.0
awesome	NaN	30.0

Many-to-one & many-to-many joins

One-to-one joins

- So far in this lecture, the joins we have worked with are called **one-to-one** joins.
- Neither the left DataFrame (`baby_small`) nor the right DataFrame (`nyt_small`) contained any duplicates in the join key.
- What if there are duplicated join keys, in one or both of the DataFrames we are merging?

```
In [13]: # Run this cell to set up the next example.
profs = pd.DataFrame(
    [['Sam', 'UCB', 5],
     ['Sam', 'UCSD', 5],
     ['Janine', 'UCSD', 8],
     ['Marina', 'UIC', 7],
     ['Justin', 'OSU', 5],
     ['Soohyun', 'UCSD', 2],
     ['Suraj', 'UCB', 2]],
    columns=['Name', 'School', 'Years']
)

schools = pd.DataFrame({
    'Abr': ['UCSD', 'UCLA', 'UCB', 'UIC'],
    'Full': ['University of California San Diego', 'University of California']
})

programs = pd.DataFrame({
    'uni': ['UCSD', 'UCSD', 'UCSD', 'UCB', 'OSU', 'OSU'],
    'dept': ['Math', 'HDSI', 'COGS', 'CS', 'Math', 'CS'],
    'grad_students': [205, 54, 281, 439, 304, 193]
})
```

Many-to-one joins

- Many-to-one joins are joins where **one** of the DataFrames contains duplicate values in the join key.

- The resulting DataFrame will preserve those duplicate entries as appropriate.

```
In [14]: dfs_side_by_side(profs, schools)
```

	Name	School	Years		Abr	Full
0	Sam	UCB	5	0	UCSD	University of California San Diego
1	Sam	UCSD	5			
2	Janine	UCSD	8	1	UCLA	University of California, Los Angeles
3	Marina	UIC	7			
4	Justin	OSU	5	2	UCB	University of California, Berkeley
5	Soohyun	UCSD	2			
6	Suraj	UCB	2	3	UIC	University of Illinois Chicago

Note that when merging `profs` and `schools`, the information from `schools` is duplicated.

- 'University of California, San Diego' appears three times.
- 'University of California, Berkeley' appears twice.

```
In [15]: profs.merge(schools, left_on='School', right_on='Abr', how='left')
```

```
Out[15]:
```

	Name	School	Years		Abr	Full
0	Sam	UCB	5	UCB	University of California, Berkeley	
1	Sam	UCSD	5	UCSD	University of California San Diego	
2	Janine	UCSD	8	UCSD	University of California San Diego	
3	Marina	UIC	7	UIC	University of Illinois Chicago	
4	Justin	OSU	5	NaN		NaN
5	Soohyun	UCSD	2	UCSD	University of California San Diego	
6	Suraj	UCB	2	UCB	University of California, Berkeley	

Many-to-many joins

Many-to-many joins are joins where both DataFrames have duplicate values in the join key.

```
In [16]: dfs_side_by_side(profs, programs)
```


	Name	School	Years		uni	dept	grad_students
0	Sam	UCB	5	0	UCSD	Math	205
1	Sam	UCSD	5	1	UCSD	HDSI	54
2	Janine	UCSD	8	2	UCSD	COGS	281
3	Marina	UIC	7	3	UCB	CS	439
4	Justin	OSU	5	4	OSU	Math	304
5	Soohyun	UCSD	2	5	OSU	CS	193
6	Suraj	UCB	2				

Before running the following cell, try predicting the number of rows in the output.

In [17]: `profs.merge(programs, left_on='School', right_on='uni')`

Out[17]:

	Name	School	Years	uni	dept	grad_students
0	Sam	UCB	5	UCB	CS	439
1	Sam	UCSD	5	UCSD	Math	205
2	Sam	UCSD	5	UCSD	HDSI	54
...
10	Soohyun	UCSD	2	UCSD	HDSI	54
11	Soohyun	UCSD	2	UCSD	COGS	281
12	Suraj	UCB	2	UCB	CS	439

13 rows x 6 columns

- `merge` stitched together every UCSD row in `profs` with every UCSD row in `programs`.
- Since there were 3 UCSD rows in `profs` and 3 in `programs`, there are $3 \times 3 = 9$ UCSD rows in the output. The same applies for all other schools.

Question 🤔

Fill in the blank so that the last statement evaluates to `True`.

```
df = profs.merge(programs, left_on='School', right_on='uni')
df.shape[0] == (____).sum()
```

Don't use `merge` (or `join`) in your solution!

```
In [18]: dfs_side_by_side(profs, programs)
```

	Name	School	Years		uni	dept	grad_students
0	Sam	UCB	5	0	UCSD	Math	205
1	Sam	UCSD	5	1	UCSD	HDSI	54
2	Janine	UCSD	8	2	UCSD	COGS	281
3	Marina	UIC	7	3	UCB	CS	439
4	Justin	OSU	5	4	OSU	Math	304
5	Soohyun	UCSD	2	5	OSU	CS	193
6	Suraj	UCB	2				

```
In [19]: # Your code goes here.
```

Returning back to our original question

Let's find the popularity of baby name categories over time. To start, we'll define a DataFrame that has one row for every combination of 'category' and 'Year'.

```
In [20]: cate_counts = (
    baby
    .merge(nyt, left_on='Name', right_on='nyt_name')
    .groupby(['category', 'Year'])
    ['Count']
    .sum()
    .reset_index()
)
cate_counts
```


```
Out[20]:
```

	category	Year	Count
0	boomer	1880	292
1	boomer	1881	298
2	boomer	1882	326
...
659	mythology	2020	3516
660	mythology	2021	3895
661	mythology	2022	4049

662 rows × 4 columns

```
In [21]: # We'll talk about plotting code soon!
import plotly.express as px
```

```
fig = px.line(cate_counts, x='Year', y='Count',
              facet_col='category', facet_col_wrap=3,
              facet_row_spacing=0.15,
              width=600, height=400)
fig.update_yaxes(matches=None, showticklabels=False)
fig
```



Transforming with `.apply`

Transforming values

- A **transformation** results from performing some operation on every element in a sequence, e.g. a Series.
- While we haven't discussed it yet in DSC 80, you learned how to transform Series in DSC 10, using the `apply` method. `apply` is very flexible – it takes in a function, which itself takes in a single value as input and returns a single value.

In [22]: baby

Out [22]:

	Name	Sex	Count	Year
0	Liam	M	20456	2022
1	Noah	M	18621	2022
2	Olivia	F	16573	2022
...
2085155	Wright	M	5	1880
2085156	York	M	5	1880
2085157	Zachariah	M	5	1880

2085158 rows × 4 columns

```
In [23]: def number_of_vowels(string):
          return sum(c in 'aeiou' for c in string.lower())

          baby['Name'].apply(number_of_vowels)
```

```
Out [23]: 0          2
          1          2
          2          4
          ..
          2085155    1
          2085156    1
          2085157    4
          Name: Name, Length: 2085158, dtype: int64
```

```
In [24]: # Built-in functions work with apply, too.
          baby['Name'].apply(len)
```

```
Out [24]: 0          4
          1          4
          2          6
          ..
          2085155    6
          2085156    4
          2085157    9
          Name: Name, Length: 2085158, dtype: int64
```

The price of `apply`

Unfortunately, `apply` runs really slowly!

```
In [25]: %%timeit
          baby['Name'].apply(number_of_vowels)
```

1.1 s ± 22.3 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
In [ ]: %%timeit
          res = []
```

```
for name in baby['Name']:
    res.append(number_of_vowels(name))
```

Internally, `apply` actually just runs a `for` -loop in Python!

So, when possible – say, when applying arithmetic operations – we should work on Series objects directly and avoid `apply` !

The price of `apply`

```
In [ ]: %%timeit
        baby['Year'].apply(lambda y: y // 10 * 10)
```

```
In [ ]: %%timeit
        baby['Year'] // 10 * 10 # Rounds down to the nearest multiple of 10.
```

Apply is 100x slower!

The `.str` accessor

For string operations, `pandas` provides a convenient `.str` accessor.

```
In [ ]: %%timeit
        baby['Name'].str.upper()
```

```
In [ ]: %%timeit
        baby['Name'].apply(lambda s: s.upper())
```

It's very convenient and **runs about the same speed as `apply` !**

Dataset overview

San Diego food safety

From [this article](#) ([archive link](#)):

In the last three years, one third of San Diego County restaurants have had at least one major food safety violation.

99% Of San Diego Restaurants Earn 'A' Grades, Bringing Usefulness of System Into Question

From [this article](#) ([archive link](#)):

Food held at unsafe temperatures. Employees not washing their hands.
Dirty countertops. Vermin in the kitchen. An expired restaurant permit.

Restaurant inspectors for San Diego County found these violations during
a routine health inspection of a diner in La Mesa in November 2016.

Despite the violations, the restaurant was awarded a score of 90 out of
100, the lowest possible score to achieve an 'A' grade.

The data

- We downloaded the data about the 1000 restaurants closest to UCSD from [here](#).
- We had to download the data as JSON files, then process it into DataFrames. You'll learn how to do this soon!
 - Until now, you've (largely) been presented with CSV files that `pd.read_csv` could load without any issues.
 - But there are many different formats and possible issues when loading data in from files.
 - See [Chapter 8 of Learning DS](#) for more.

```
In [ ]: rest_path = Path('data') / 'restaurants.csv'
        insp_path = Path('data') / 'inspections.csv'
        viol_path = Path('data') / 'violations.csv'
```

```
In [ ]: rest = pd.read_csv(rest_path)
        insp = pd.read_csv(insp_path)
        viol = pd.read_csv(viol_path)
```

Question 🤔

The first article said that one third of restaurants had at least one major safety violation. Which DataFrames and columns seem most useful to verify this?

```
In [ ]: rest.head(2)
```

```
In [ ]: rest.columns
```

```
In [ ]: insp.head(2)
```

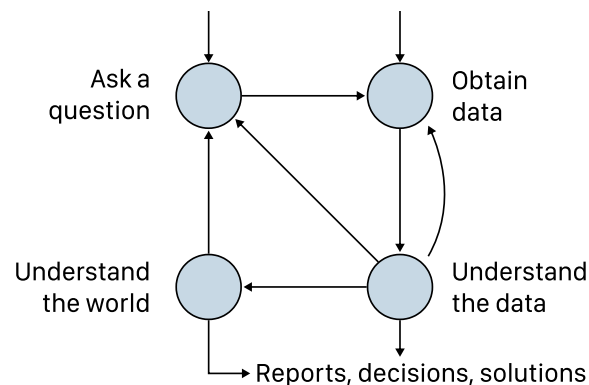
```
In [ ]: insp.columns
```

```
In [ ]: viol.head(2)
```

```
In [ ]: viol.columns
```

Exploratory data analysis and feature types

The data science lifecycle, revisited



We're at the stage of **understanding the data**.

Exploratory data analysis (EDA)

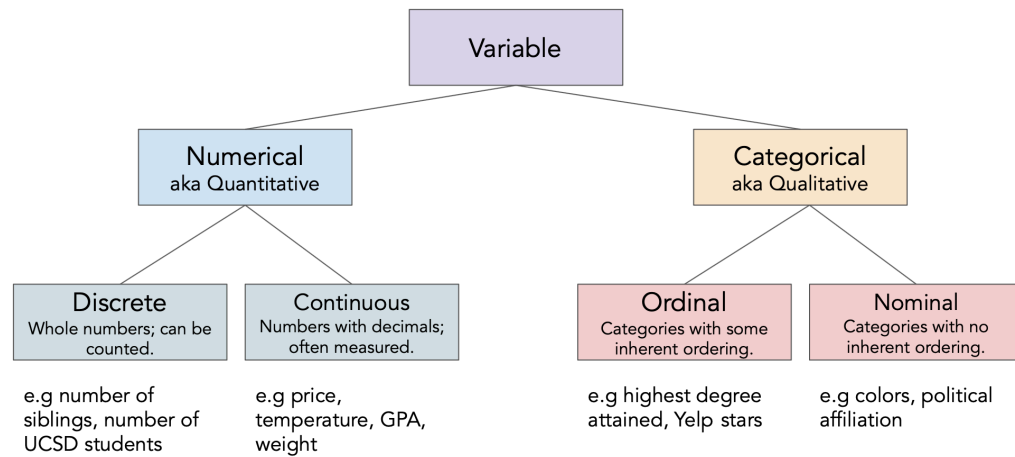
- Historically, data analysis was dominated by formal statistics, including tools like confidence intervals, hypothesis tests, and statistical modeling.
- In 1977, John Tukey [defined](#) the term **exploratory data analysis**, which described a philosophy for proceeding about data analysis:

Exploratory data analysis is actively incisive, rather than passively descriptive, with real emphasis on the discovery of the unexpected.

- Practically, EDA involves, among other things, computing summary statistics and drawing plots to understand the nature of the data at hand.

The greatest gains from data come from surprises... The unexpected is best brought to our attention by **pictures**.

Different feature types



Note that numerical variables can be stored as strings, and categorical variables can be stored as numbers!

Question 🤔

Determine the **feature type** of each of the following variables.

- `insp['score']`
- `insp['grade']`
- `viol['violation_accela']`
- `viol['major_violation']`
- `rest['business_id']`
- `rest['opened_date']`

In []: *# Your code goes here.*

Feature types vs. data types

- The data type `pandas` uses is not the same as the "data type" we talked about just now!
 - There's a difference between feature type (which has to do with the data's meaning) and computational data type (which has to do with how it is stored).
- Take care when the two don't match up very well!

In []: *# pandas stores these as ints, but they're actually nominal.*
`rest['business_id']`

In []: *# pandas stores these as strings, but they're actually numeric.*
`rest['opened_date']`

Data cleaning

Four pillars of data cleaning

When loading in a dataset, to clean the data – that is, to prepare it for further analysis – we will:

1. Perform **data quality checks**.
2. Identify and handle **missing values**.
3. Perform **transformations**, including converting time series data to **timestamps**.
4. Modify **structure** as necessary.

Data cleaning: Data quality checks

Data quality checks

We often start an analysis by checking the quality of the data.

- Scope: Do the data match your understanding of the population?
- Measurements and values: Are the values reasonable?
- Relationships: Are related features in agreement?
- Analysis: Which features might be useful in a future analysis?

Scope

Do the data match your understanding of the population?

We were told that we're only looking at the 1000 restaurants closest to UCSD, so the restaurants in `rest` should agree with that.

```
In [ ]: rest.sample(5)
```

Measurements and values

Are the values reasonable?

Do the values in the `'grade'` column match what we'd expect grades to look like?

```
In [ ]: insp['grade'].value_counts()
```

What kinds of information does the `insp` DataFrame hold?

```
In [ ]: insp.info()
```

What's going on in the `'address'` column of `rest` ?

```
In [ ]: # Are there multiple restaurants with the same address?
rest['address'].value_counts()
```

```
In [ ]: # Keeps all rows with duplicate addresses.
(
    rest
    .groupby('address')
    .filter(lambda df: df.shape[0] >= 2)
    .sort_values('address')
)
```

```
In [ ]: # Does the same thing as above!
(
    rest[rest.duplicated(subset=['address'], keep=False)]
    .sort_values('address')
)
```

Relationships

Are related features in agreement?

Do the `'address'` es and `'zip'` codes in `rest` match?

```
In [ ]: rest[['address', 'zip']]
```

What about the `'score'` s and `'grade'` s in `insp` ?

```
In [ ]: insp[['score', 'grade']]
```

Analysis

Which features might be useful in a future analysis?

- We're most interested in:
 - These columns in the `rest` DataFrame: `'business_id'`, `'name'`, `'address'`, `'zip'`, and `'opened_date'`.
 - These columns in the `insp` DataFrame: `'business_id'`, `'inspection_id'`, `'score'`, `'grade'`, `'completed_date'`, and `'status'`.
 - These columns in the `viol` DataFrame: `'inspection_id'`, `'violation'`, `'major_violation'`, `'violation_text'`, and

`'violation_accela'.`

- Also, let's rename a few columns to make them easier to work with.

💡 Pro-Tip: Using `pipe`

When we manipulate DataFrames, it's best to define individual functions for each step, then use the `pipe` method to chain them all together.

The `pipe` DataFrame method takes in a function, which itself takes in a DataFrame and returns a DataFrame.

- In practice, we would add functions one by one to the top of a notebook, then `pipe` them all.
- For today, will keep re-running `pipe` to show data cleaning process.

```
In [ ]: def subset_rest(rest):
        return rest[['business_id', 'name', 'address', 'zip', 'opened_date']]

rest = (
    pd.read_csv(rest_path)
    .pipe(subset_rest)
)
rest
```

```
In [ ]: # Same as the above – but the above makes it easier to chain more .pipe calls
subset_rest(pd.read_csv(rest_path))
```

Let's use `pipe` to keep (and rename) the subset of the columns we care about in the other two DataFrames as well.

```
In [ ]: def subset_insp(insp):
        return (
            insp[['business_id', 'inspection_id', 'score', 'grade', 'completed_date']]
            .rename(columns={'completed_date': 'date'})
        )

insp = (
    pd.read_csv(insp_path)
    .pipe(subset_insp)
)
```

```
In [ ]: def subset_viol(viol):
        return (
            viol[['inspection_id', 'violation', 'major_violation', 'violation_acceleration']]
            .rename(columns={'violation': 'kind',
                              'major_violation': 'is_major',
                              'violation_acceleration': 'violation'})
        )

viol = (
```

```
pd.read_csv(viol_path)
    .pipe(subset_viol)
)
```

Combining the restaurant data

Let's join all three DataFrames together so that we have all the data in a single DataFrame.

```
In [ ]: def merge_all_restaurant_data():
        return (
            rest
            .merge(insp, on='business_id', how='left')
            .merge(viol, on='inspection_id', how='left')
        )

df = merge_all_restaurant_data()
df
```

Question 🤔

Why should the function above use two left joins? What would go wrong if we used other kinds of joins?

We will probably end lecture here.

Data cleaning: Missing values

Missing values

Next, it's important to check for and handle missing values, as they can have a big effect on your analysis.

```
In [ ]: insp[['score', 'grade']]
```

```
In [ ]: # The proportion of values in each column that are missing.
        insp.isna().mean()
```

```
In [ ]: # Why are there null values here?
        # insp['inspection_id'] and viol['inspection_id'] don't have any null values
        df[df['inspection_id'].isna()]
```

There are many ways of handling missing values, which we'll cover in an entire lecture next week. But a good first step is to check how many there are!

Data cleaning: Transformations and timestamps

Transformations and timestamps

From last class:

A transformation results from performing some operation on every element in a sequence, e.g. a Series.

It's often useful to look at ways of transforming your data to make it easier to work with.

- Type conversions (e.g. changing the string "\$2.99" to the number 2.99).
- Unit conversion (e.g. feet to meters).
- Extraction (Getting 'vermin' out of 'Vermin Violation Recorded on 10/10/2023').

Creating timestamps

Most commonly, we'll parse dates into `pd.Timestamp` objects.

```
In [ ]: # Look at the dtype!
insp['date']
```

```
In [ ]: # This magical string tells Python what format the date is in.
# For more info: https://docs.python.org/3/library/datetime.html#strftime-ar
date_format = '%Y-%m-%d'
pd.to_datetime(insp['date'], format=date_format)
```

```
In [ ]: # Another advantage of defining functions is that we can reuse this function
# for the 'opened_date' column in `rest` if we wanted to.
def parse_dates(insp, col):
    date_format = '%Y-%m-%d'
    dates = pd.to_datetime(insp[col], format=date_format)
    return insp.assign(**{col: dates})

insp = (
    pd.read_csv(insp_path)
    .pipe(subset_insp)
    .pipe(parse_dates, 'date')
)

# We should also remake df, since it depends on insp.
```

```
# Note that the new insp is used to create df!
df = merge_all_restaurant_data()
```

```
In [ ]: # Look at the dtype now!
df['date']
```

Working with timestamps

- We often want to adjust granularity of timestamps to see overall trends, or seasonality.
- Use the `resample` method in `pandas` ([documentation](#)).
 - Think of it like a version of `groupby`, but for timestamps.
 - For instance, `insp.resample('2W', on='date')` separates every two weeks of data into a different group.

```
In [ ]: insp.resample('2W', on='date')['score'].mean()
```

```
In [ ]: # Where are those numbers coming from?
insp[
    (insp['date'] >= pd.Timestamp('2020-01-05')) &
    (insp['date'] < pd.Timestamp('2020-01-19'))
]['score']
```

```
In [ ]: (insp.resample('2W', on='date')
        .size()
        .plot(title='Number of Inspections Over Time')
        )
```

The `.dt` accessor

Like with Series of strings, `pandas` has a `.dt` accessor for properties of timestamps ([documentation](#)).

```
In [ ]: insp['date']
```

```
In [ ]: insp['date'].dt.day
```

```
In [ ]: insp['date'].dt.dayofweek
```

```
In [ ]: dow_counts = insp['date'].dt.dayofweek.value_counts()
fig = px.bar(dow_counts)
fig.update_xaxes(tickvals=np.arange(7), ticktext=['Mon', 'Tues', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun'])
```

Data cleaning: Modifying structure

Reshaping DataFrames

We often **reshape** the DataFrame's structure to make it more convenient for analysis. For example, we can:

- Simplify structure by removing columns or taking a set of rows for a particular period of time or geographic area.
 - We already did this!
- Adjust granularity by aggregating rows together.
 - To do this, use `groupby` (or `resample`, if working with timestamps).
- Reshape structure, most commonly by using the DataFrame `melt` method to unpivot a dataframe.

Using `melt`

- The `melt` method is common enough that we'll give it a special mention.
- We'll often encounter pivot tables (esp. from government data), which we call *wide* data.
- The methods we've introduced work better with *long-form* data, or *tidy* data.
- To go from wide to long, `melt`.

Wide-form:

	Jan	Feb	Mar
2001	10	20	30
2002	130	200	340

Long-form:

Year	Month	Seats
2001	Jan	10
2001	Feb	20
2001	Mar	30
2002	Jan	130
2002	Feb	200
2002	Mar	340

Example usage of `melt`

```
In [ ]: wide_example = pd.DataFrame({
    'Year': [2001, 2002],
    'Jan': [10, 130],
    'Feb': [20, 200],
    'Mar': [30, 340]
}).set_index('Year')
wide_example
```

```
In [ ]: wide_example.melt(ignore_index=False)
```

Exploration

Question 🤔 (Answer at dsc80.com/q)

Code: `qs`

What questions do you want me to try and answer with the data? I'll start with a single pre-prepared question, and then answer student questions until we run out of time.

Example question: Can we rank restaurants by their number of violations? How about separately for each zip code?

And why would we want to do that? 🤔

In []:

Summary, next time

Summary

- Data cleaning is a necessary starting step in data analysis. There are four pillars of data cleaning:
 - Quality checks.
 - Missing values.
 - Transformations and timestamps.
 - Modifying structure.
- Approach EDA with an open mind, and draw lots of visualizations.

Next time

Hypothesis and permutation testing. Some of this will be DSC 10 review, but we'll also push further!